The Performance Commons — List of Figures and Tables

Read top to bottom, the visuals carry the paper's argument on their own: magnitude → mechanism → the same failure everywhere → the same problem, distributed → proof → destination → already underway → the program. Each entry notes the beat it carries and whether it rests on measured data or is an argued/conceptual construct. (The defections and the diagnosis of the herd live in §2; Figure 2 carries §2.2's measured centerpiece.)

---

latest progress and open items for each, stated plainly. *Beat: the questions are under active attack, and the breadth of the program is itself evidence of the §6 multiplier.* — **status report, self-reported**; states verified against the repos at the time of writing.

Table A2 — The wedge economics. (Appendix A) macOS CI and inference prices, mid-2026: GitHub macOS runner $3.72/hr (10× Linux) and EC2 Mac $0.65/hr (24-hr minimum) vs ~$0.006/hr idle-Mac electricity; 70B unified-memory inference ~$0.60/Mtoken vs $0.40 cheapest API and ~$0.06 batched-4090 small models. *Beat: the Apple wedge prices where licensing and unified memory are the moat.* — **quantitative** (vendor price lists, published benchmarks); spot rates fluctuate.

How the sequence maps to the prose

Visual	Section	Argument beat
Table 1	§1.1	the gap is real and large
Figure 1	§1.2	the gap is effort, not physics (Q1: hardware)
—	§2	privilege → the central claim: continuous optimization is an under-provided public good (prose only)
Figure 2	§2.2	the embargo natural experiment: chips restricted, parity anyway
Table 2	§3	the same failure is everywhere in the stack
Figure 3	§4	distributed computing is the same problem — the wasted supply (Q2: software)
Figure 4	§4	unbounded machine demand makes the idle resource valuable
Table 3	§5	the entry cost of first silicon has collapsed
Figure 5	§5	the commons survives its platform's death, measured
—	§6–§7	(AI multiplies the commons; the vendor is the aggregation point — prose only)
Figure 6	§8	where the optimum lands as the commons matures
Figure 7	§8	and it is already underway
—	§9	honest limits (prose only)
—	Conclusion	build collaboration-native hardware (prose only)
Table A1	Appendix A	the program in detail, honestly
Table A2	Appendix A	the wedge economics, priced

The "needs a bigger project" items (flagged for honesty)

- **Figure 6's boundary** — the real position of the ASIC / CPU / CGRA optimal regions needs the whole-system workload characterization §9 calls the prerequisite (achieved-vs-achievable parallelism, dataflow locality, phase and input statistics across the real consumer mix).
- **Figure 4's demand curve** — conceptual; the ~$0.50 vs ~$0.08 cost anchors are from Omerta's reliability-market simulation and the task-value ladder from its economic-analysis notes. Omerta itself is partially built: a substantial, well-tested mesh-networking layer (encryption, NAT traversal, gossip discovery, tunnelling, a mesh-native virtual network, plus Terraform-managed bootstrap/STUN servers) and working ephemeral-VM and simulation components, but the end-to-end compute loop is broken mid-migration and the trust/payment chain is validated in simulation rather than running live (two of six transaction types tested).
- **Figure 7's share curve and projection** — the adoption trend is illustrative; the dated milestones and the $16.7B / $49B acquisition values are real.
- **Figure 1's finite-effort curve** — illustrative; only the realizable-frontier endpoints (the intrinsic fabric tax) are measured.